

Kaggle NLP Approaches with discussion on Kaggle Text Classification Challenge

| Date : 15th June , 2019 | Venue : Grofers , Gurugram | Speaker : Vikas Kumar |

About Myself

- Vikas Kumar
Consultant @ Deloitte USI
Analytics & Cognitive, Applied AI Team
(since April 2019)

Background

- Consultant @EXL
- Graduation from IIT Kharagpur (2017)
- Working on NLP and Computer Vision problems
- Active Kagglers since 2017 (3 Silver medals , 1 Bronze)

Linkedin: [linkedin.com/in/vikashiitkgp/](https://www.linkedin.com/in/vikashiitkgp/)

Kaggle: [kaggle.com/vikas15](https://www.kaggle.com/vikas15)



Competitions Expert		
Current Rank	Highest Rank	
1976	1499	
of 112,358		
		
0	3	1
2018 Data Science Bowl		172 nd
 a year ago-Top 5%		of 3634
Quora Insincere Questions ...		173 rd
 4 months ago-Top 5%		of 4037
TalkingData AdTracking Fr...		193 rd
 a year ago-Top 5%		of 3951

The Audience

- Know some Analytics already?
- Familiar with Natural Language Processing?
- Expertise in deep learning?

Use cases of NLP

Information Retrieval

Doc A

Doc 1

Doc 2

Doc 3

Sentiment Analysis



Information Extraction



Machine Translation



Natural Language Processing

Question Answering

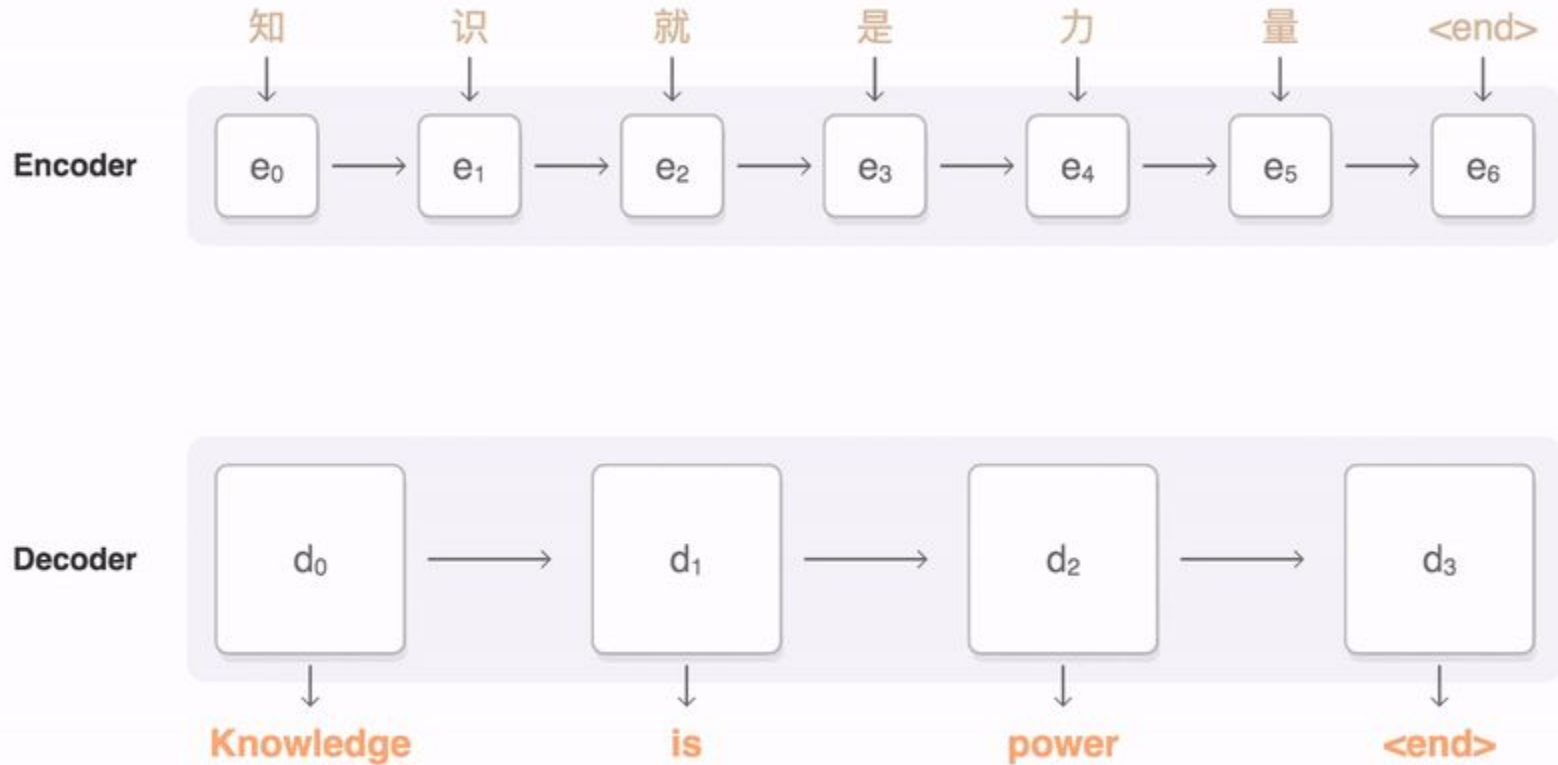


Human: When was Apollo sent to space?



Machine: First flight - AS-201, February 26, 1966

Machine Translation



Name Entity Recognition

"There was nothing about this storm that was as expected," said **Jeff Masters**, a meteorologist and founder of **Weather Underground**. "**Irma** could have been so much worse. If it had traveled 20 miles north of the coast of **Cuba**, you'd have been looking at a (Category) 5 instead of a (Category) 3."

Person

Organization

Location

Question Answer System

The Normans (Norman: Nourmands; French: Normands; Latin: Normanni) were the people who in the 10th and 11th centuries gave their name to Normandy, a region in France. They were descended from Norse ("Norman" comes from "Norseman") raiders and pirates from Denmark, Iceland and Norway who, under their leader Rollo, agreed to swear fealty to King Charles III of West Francia. Through generations of assimilation and mixing with the native Frankish and Roman-Gaulish populations, their descendants would gradually merge with the Carolingian-based cultures of West Francia. The distinct cultural and ethnic identity of the Normans emerged initially in the first half of the 10th century, and it continued to evolve over the succeeding centuries.

When were the Normans in Normandy?

- *Ground Truth Answers:* 10th and 11th centuries
Prediction: 10th and 11th centuries

In what country is Normandy located?


- *Ground Truth Answers:* France
Prediction: France

Text Classification

Spam Filter



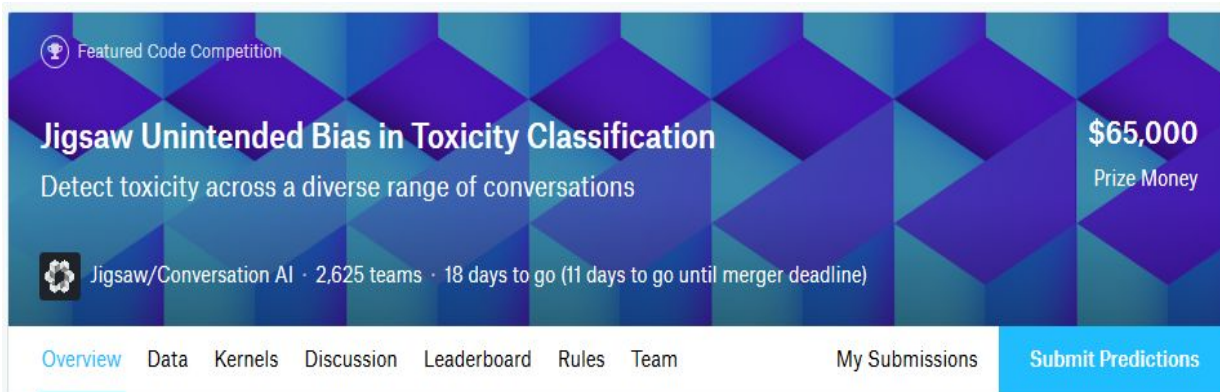
Text Classification Example



Kaggle Text Classification Challenge



Kaggle Toxic Classification Challenge



Featured Code Competition

Jigsaw Unintended Bias in Toxicity Classification

Detect toxicity across a diverse range of conversations

\$65,000
Prize Money

Jigsaw/Conversation AI · 2,625 teams · 18 days to go (11 days to go until merger deadline)

[Overview](#) [Data](#) [Kernels](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#) [My Submissions](#) [Submit Predictions](#)

[Competition Link](#)

Problem Statement :-

- To detect toxic comments and minimize unintended model bias
- Build toxicity models that operate fairly across a diverse range of conversations

Last day of submission - 27 June, 2019

Sample Dataset

comment_text	target
And, as an American, I am happy to tell you freeloaders to go pound sand.	0
And they are receiving all the perks of being Canadian without paying into our taxes.	0
I'm embarrassed to be called an American with this buffoon leading the country.	1
Just remember, men wore jocks long before they wore helmets. Priorities!	0
good points mister gay, our mayor spendowitz better get on the job eh.	1
If motley was trolling then i could have a problem with his posts and anonymity. As i see it he doesnt troll and see no reason to know his backgroud information. Paul didnt you try and link with me on linkedin. If i wanted people to dig up personal information about me i would just give it to them.	0
Anon	
The royal parasites are back!	1
and you have to say that about the liberals too idiot, more voted against justin than voted for him! give it a rest, this bs is meaningless	1
All form of hate need to be eliminated racism, homophobia, Islamophobia, Anti-Semitism, discrimination against Transgenders. We must get behind Justin, Obama, Hillary and Premier Wynne and fight Trump and this racist Quebec law!	0
Dont do business with Wells Fargo,period, divest.	0
THAT IS A STUPID COMMENT,COULD YOU HAVE DONE BETTER.	1

Approach

What top kagglers are using for Text classification ?

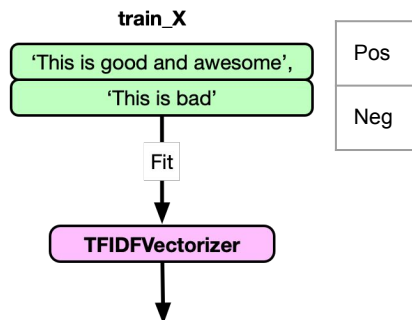
Baseline Models: Text Classification

2012	Bag of words (tf-idf)
2013	LSTM (word Embeddings - word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFiT, BERT) (Transfer Learning in NLP)

Traditional Approach : TF-IDF

2012	Bag of words (tf-idf)
2013	LSTM (word Embeddings - word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFit, BERT) (Transfer Learning in NLP)

Traditional Approach : TF-IDF



TF-IDF = Term Frequency x Inverse Document Frequency

$TF(t) = (\text{Number of times term } t \text{ appears in a document}) / (\text{Total number of terms in the document})$

$IDF(t) = \log (\text{Total number of documents} / \text{Number of documents with term } t \text{ in it})$

Term Frequency

No of time word appears/No of total terms in Document

This	is	good	bad	awesome	and
1/5	1/5	1/5	0	1/5	1/5
1/3	1/3	0	1/3	0	0

Inverse Document Frequency

$-\log(\text{ratio of documents that include the word})$

This	is	good	bad	awesome	and
$\log(2/2)$	$\log(2/2)$	$\log(2/1)$	$\log(2/1)$	$\log(2/1)$	$\log(2/1)$

X

Features

This	is	good	bad	awesome	and
0	0	$1/5 * \log(2/1)$	0	$1/5 * \log(2/1)$	$1/5 * \log(2/1)$
0	0	0	$1/3 * \log(2/1)$	0	0

Pos

Neg

[Reference](#)

TF-IDF Shortcoming

Is sequence of words in sentence important?

He only could solve that problem. —————→ Positive

He could solve only that problem. —————→ Negative


RNN / LSTM (word Embeddings - word2vec)

2012	Bag of words (tf-idf)
2013	RNN / LSTM (word Embeddings - word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFit, BERT) (Transfer Learning in NLP)

RNN - Recurrent Neural Network


LSTM - Long Short Term Memory

RNN




[Reference](#)

LSTM



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$h_t = o_t * \tanh(C_t)$$




$$f_t = \sigma(W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$


$$o_t = \sigma(W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$

Reference

Word Embeddings



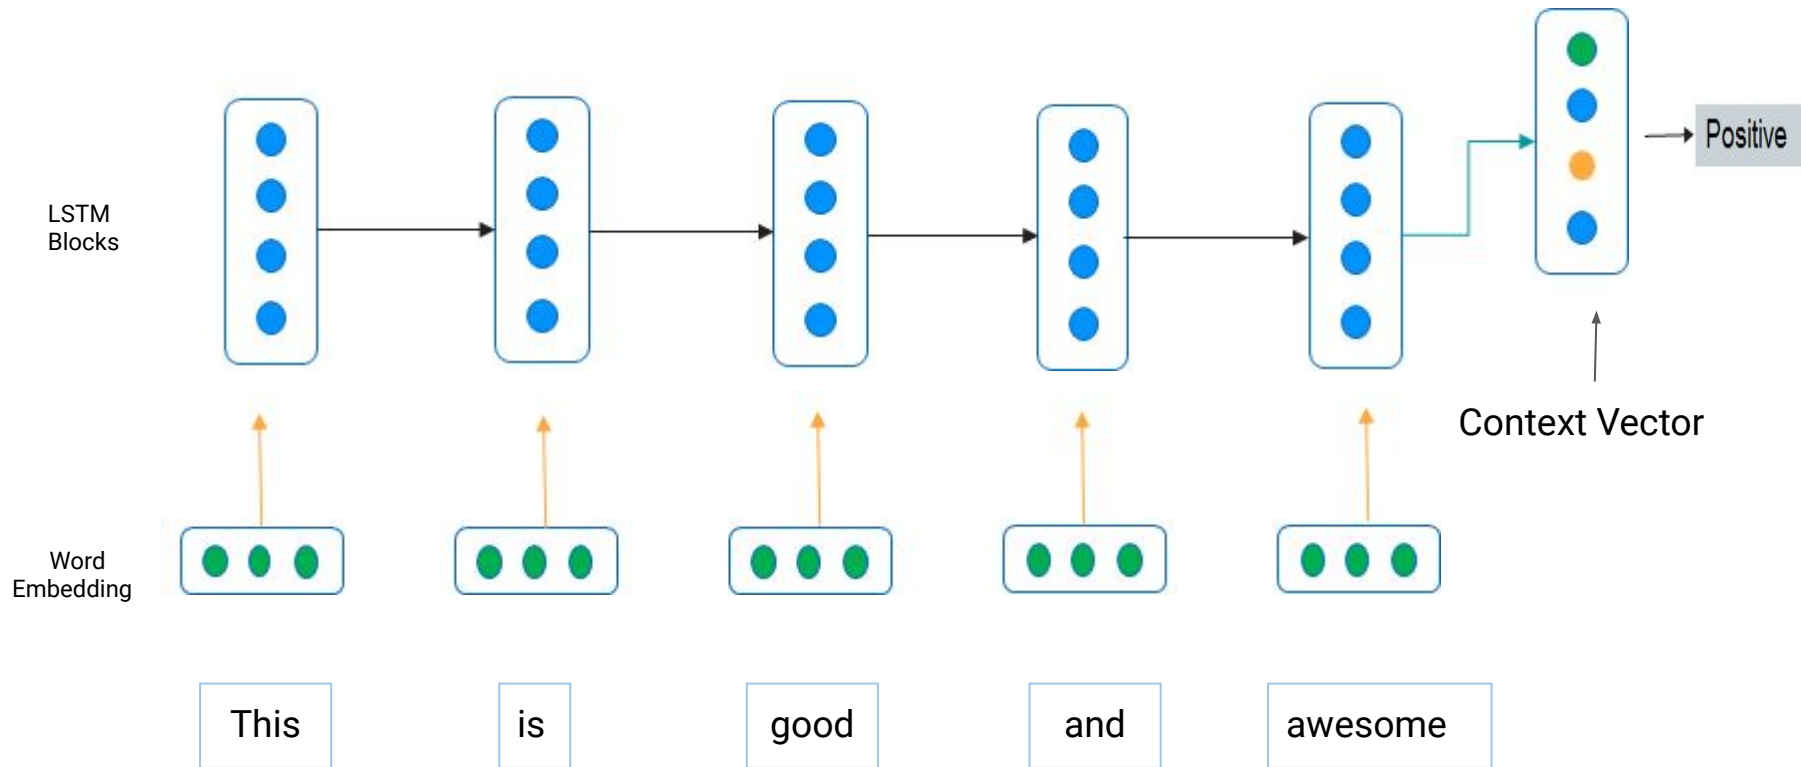
Word2Vec



Other Embeddings:

- [Pennington et al. \(2014\)](#): GloVe
- [Bojanowski et al. \(2017\)](#): fastText

LSTM Text Classification



LSTM (word Embeddings + attention)

2012	Bag of words (tf-idf)
2013	LSTM (word Embeddings - word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFit, BERT) (Transfer Learning in NLP)


Why Attention?

One of his best and most shocking films ! divine is the most filthy person ever ! mink stole also delivers a superb performance

This is terrible movie , don't waste your money on it, don't even watch it for free , that is what i have to say.

One of his best and most shocking films ! divine is the most filthy person ever ! mink stole also delivers a superb performance

Attention



$$u_i = \tanh(W_s h_i + b_s),$$

$$\alpha_i = \frac{\exp(u_i^\top u_s)}{\sum_i \exp(u_i^\top u_s)},$$

$$v = \sum_i \alpha_i h_i,$$

Language Model

2012	Bag of words (tf-idf)
2013	LSTM (word Embeddings- word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFit, BERT) (Transfer Learning in NLP)

Language Model: Frequency based

Chain Rule:

- $P(A,B,C,D) = P(A)P(B|A)P(C|A,B)P(D|A,B,C)$
- $p(B|A) = P(A,B)/P(A)$

$$P(\text{"its water is so transparent"}) = P(\text{its}) \times P(\text{water}|\text{its}) \times P(\text{is}|\text{its water}) \times P(\text{so}|\text{its water is}) \times P(\text{transparent}|\text{its water is so})$$

Bi-gram Model

Library - **kenlm**

$$P(\text{"its water is so transparent"}) = P(\text{its}) \times P(\text{water}|\text{its}) \times P(\text{is}|\text{water}) \times P(\text{so}|\text{is}) \times P(\text{transparent}|\text{so})$$

Spell - correction Example

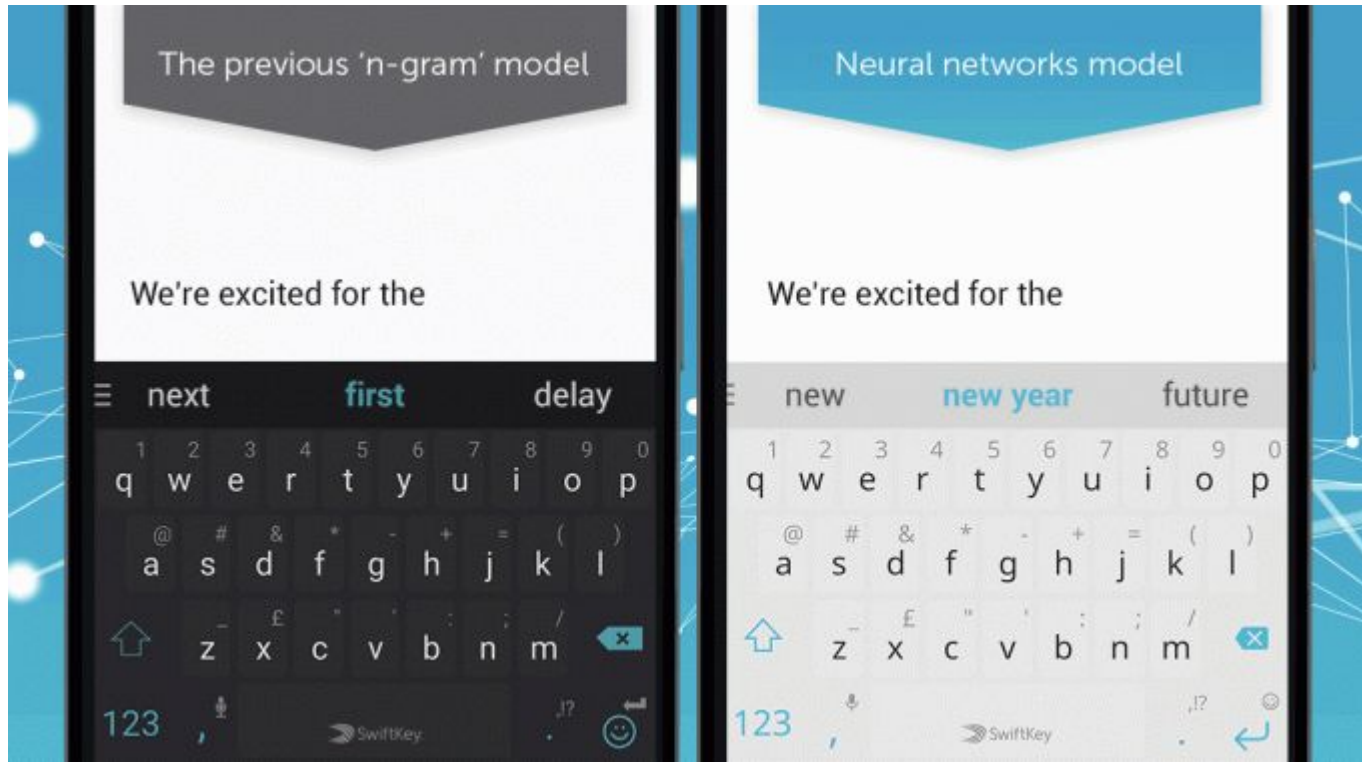
versatile **actress** whose

$P(\text{"versatile actress whose"})$


>

$P(\text{"versatile across whose"})$

SwiftKey



Neural Language Model



- generate most likely output depending on input

RNN / LSTM Block

Language Model: Text Generation

GPT2 Language Model



Give the model a prompt, or start telling a story.

A



[Reference](#)

Deep Learning Frameworks



TensorFlow

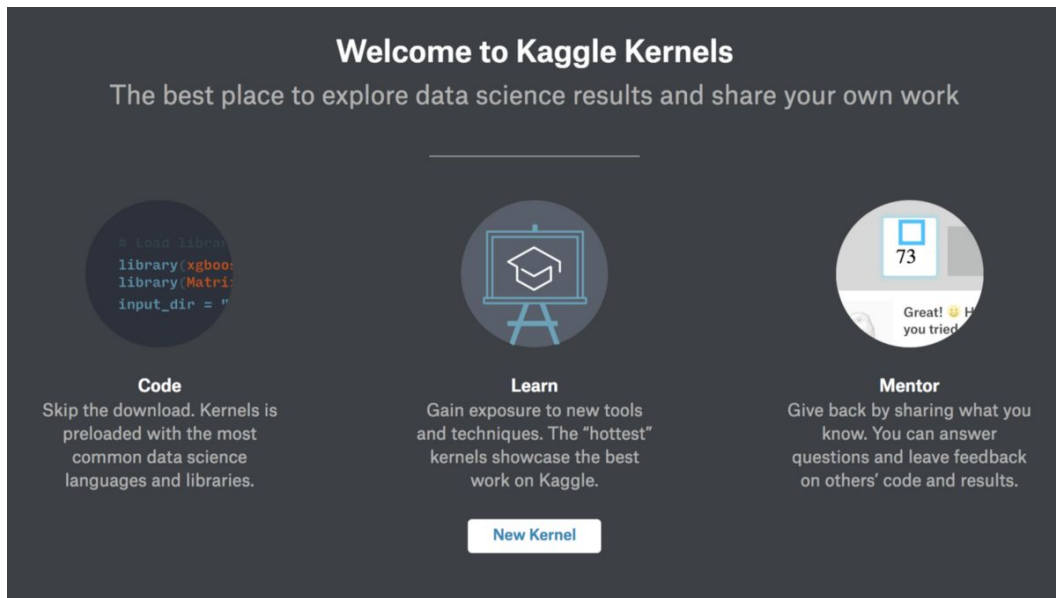


Keras

PYTORCH

fast.ai

Kaggle Kernels



[Documentation](#)

- a free cloud based platform to run code in the browser (Scripts / Jupyter Notebooks)
- can write scripts in R as well as Python
- Install any package and use any dataset for experiment
- Free GPU (~ 10 hrs for a kernel)
- can run more than 8 kernels in parallel

Transfer Learning (ULMFiT, BERT)

2012	Bag of words (tf-idf)
2013	LSTM (word Embeddings- word2vec)
2015	LSTM (word Embeddings + attention)
2018-19	Pre-trained Language Model (ULMFiT, BERT) (Transfer Learning in NLP)

Transfer Learning

ImageNet Challenge

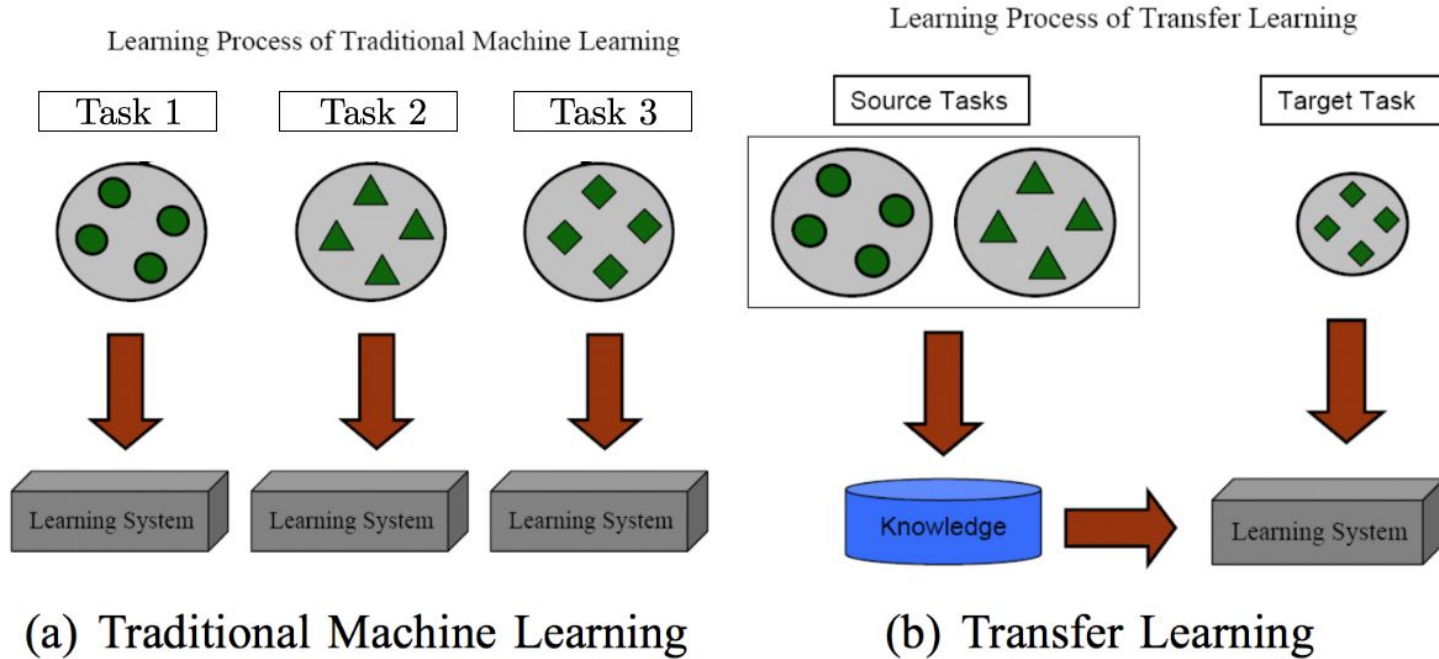
IMAGENET

- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



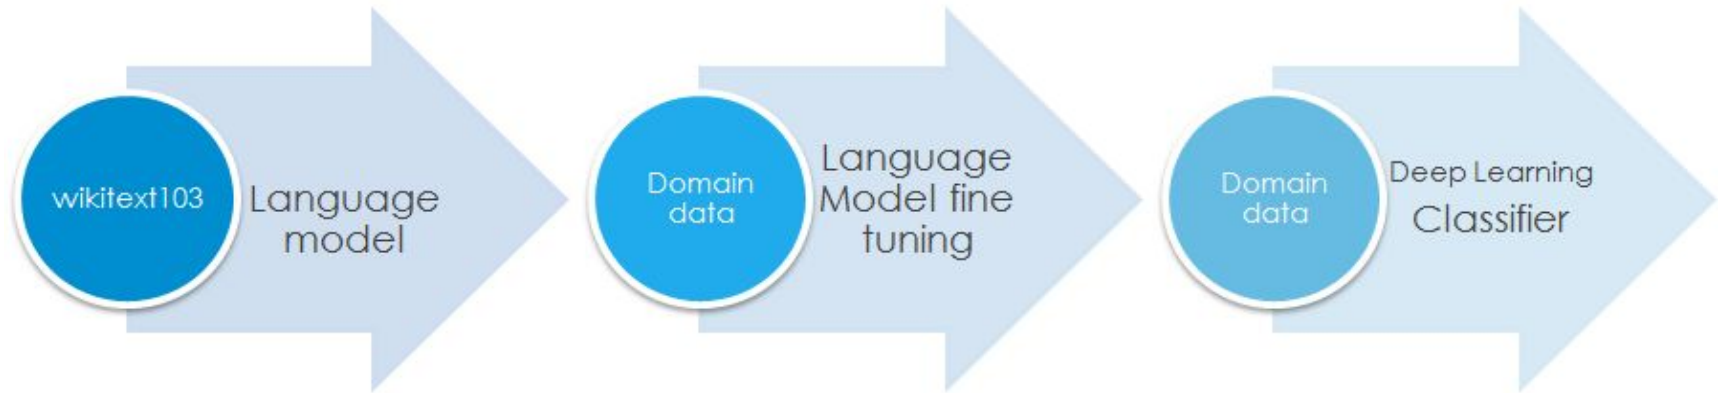
Where does Transfer Learning work ?

Transfer Learning




Transfer Learning: NLP

ULMFIT: Universal Language Model Fine tuning



[Reference](#)

ULMFiT Accuracy



SOTA results for various text classification dataset

BERT: ImageNet in NLP

Bidirectional Encoder Representations from Transformers

1 - **Semi-supervised** training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.

Semi-supervised Learning Step

Model:



Dataset:



Objective:

Predict the masked word
(language modeling)

2 - **Supervised** training on a specific task with a labeled dataset.

Supervised Learning Step

Model:
(pre-trained
in step #1)



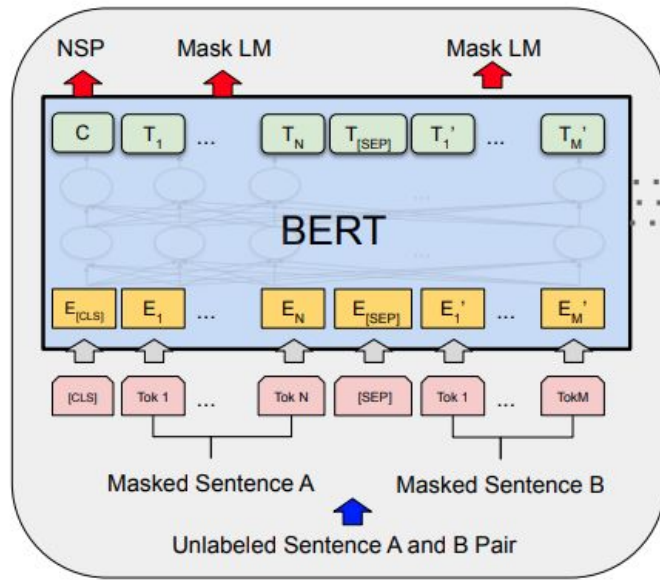
Classifier

75% Spam
25% Not Spam

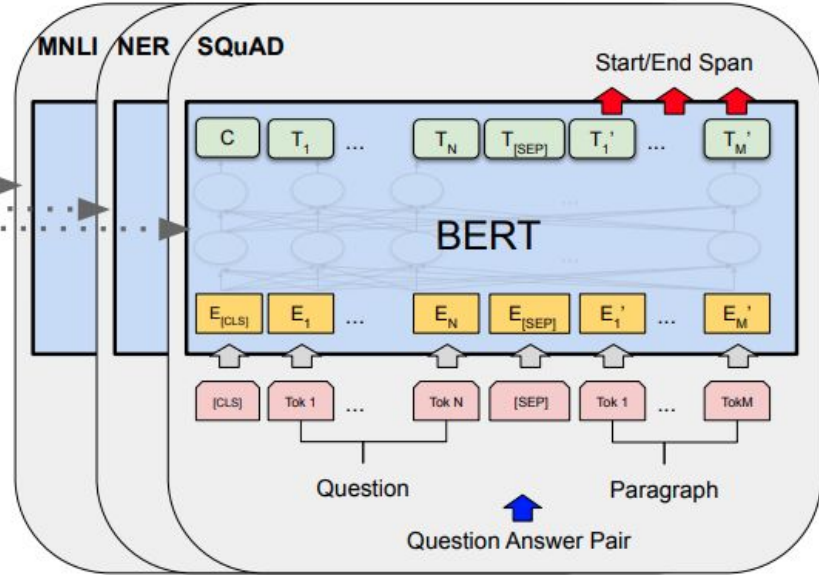
Dataset:

Email message	Class
Buy these pills	Spam
Win cash prizes	Spam
Dear Mr. Atreides, please find attached...	Not Spam

BERT: Multi-Task learning



Pre-training



Fine-Tuning

SOTA for different tasks like Text classification, NER, Question - Answer

How to get started in NLP ?

Resources

Courses

- [Andrew Ng's Deep Learning Specialization](#)
- [Deep Learning for NLP with Pytorch](#)
- [CS224n: Natural Language Processing with Deep Learning](#)
- [Fastai course](#)

Blogs / Competitions

- Kaggle Competitions
- [Sebastian Ruder blog on NLP](#)
- [Medium blog - Hugging Face](#)

Conclusion



- Best place to learn state of the art method in NLP, Computer Vision & Classical ML techniques
- Free GPU for any experiment
- Join the competition early when it launches
- Try to learn from kaggle kernels
- Create proper train validation split
- Read discussion forums carefully and try to implement some of the ideas
- Apply learnings from kaggle to your projects

Happy learnings & Never Give up !!

References

1. <https://sigmoidal.io/machine-learning-terminology-explained-top-8-must-know-concepts/>
2. <https://towardsdatascience.com/introducing-natural-language-processing-nlp-series-3-covering-essentials-of-nlp-phase-2-38839ff846de>
3. <https://hackernoon.com/beam-search-attention-for-text-summarization-made-easy-tutorial-53b7186df7086>
4. <https://rajpurkar.github.io/>
5. https://mlwhiz.com/blog/2019/01/17/deeplearning_nlp_preprocess/
6. <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>
7. <https://nlp.stanford.edu/projects/glove/>
8. <http://nlp.fast.ai/classification/2018/05/15/introducing-ulmfit.html>
9. <http://runder.io/>
10. <https://venturebeat.com/2019/02/14/openai-let-us-generate-text-with-an-ai-model-that-achieves-state-of-the-art-performance-in-several-nlp-tasks/>
11. <https://blog.swiftkey.com/swiftkey-debuts-worlds-first-smartphone-keyboard-powered-by-neural-networks/>
12. <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>
13. <https://arxiv.org/abs/1607.04606>
14. <https://www.kaggle.com/c/jigsaw-unintended-bias-in-toxicity-classification/overview>
15. <http://jalammar.github.io>

Thank you !!

Q/A Session